Welcome to DSCI 101

Introduction to Data Science

Week 13 Recap

- Decision trees
  - for classification
  - for regression
- Advantages of tree models
  - interpretability
  - non-linear relationship with interactions
- Ensemble models
  - Random Forest

Week 14 Preview

- Unsupervised learning – data driven discovery
  - No data label
  - Not focus on prediction
- Dimensionality reduction – visualize high dimensional data
  - Principal Component Analysis
- Clustering – finding subgroups in data
  - K-means

Why Dimensionality Reduction?

- Dimensionality = the columns of your tabular data

- reduce overfitting
- less noisy data
- more interpretable

- less data storage
- speed up training
- more efficient

- visualize big data
- overall data pattern
- detect outliers

PCA overview

- How to visualize the entire dataset on one scatter plot
  - tabular data with n rows and p columns
  - n observations / samples / data points
  - p variables / features / dimensions
- Human eyes are not good at “seeing” > 3 variables at once
  - a scatter plot only shows two variables
  - could just plot the “most important” two variables, but can we do better?
- Principal Component Analysis (PCA)
  - the “best” low-dimensional representation of data

PCA illustration

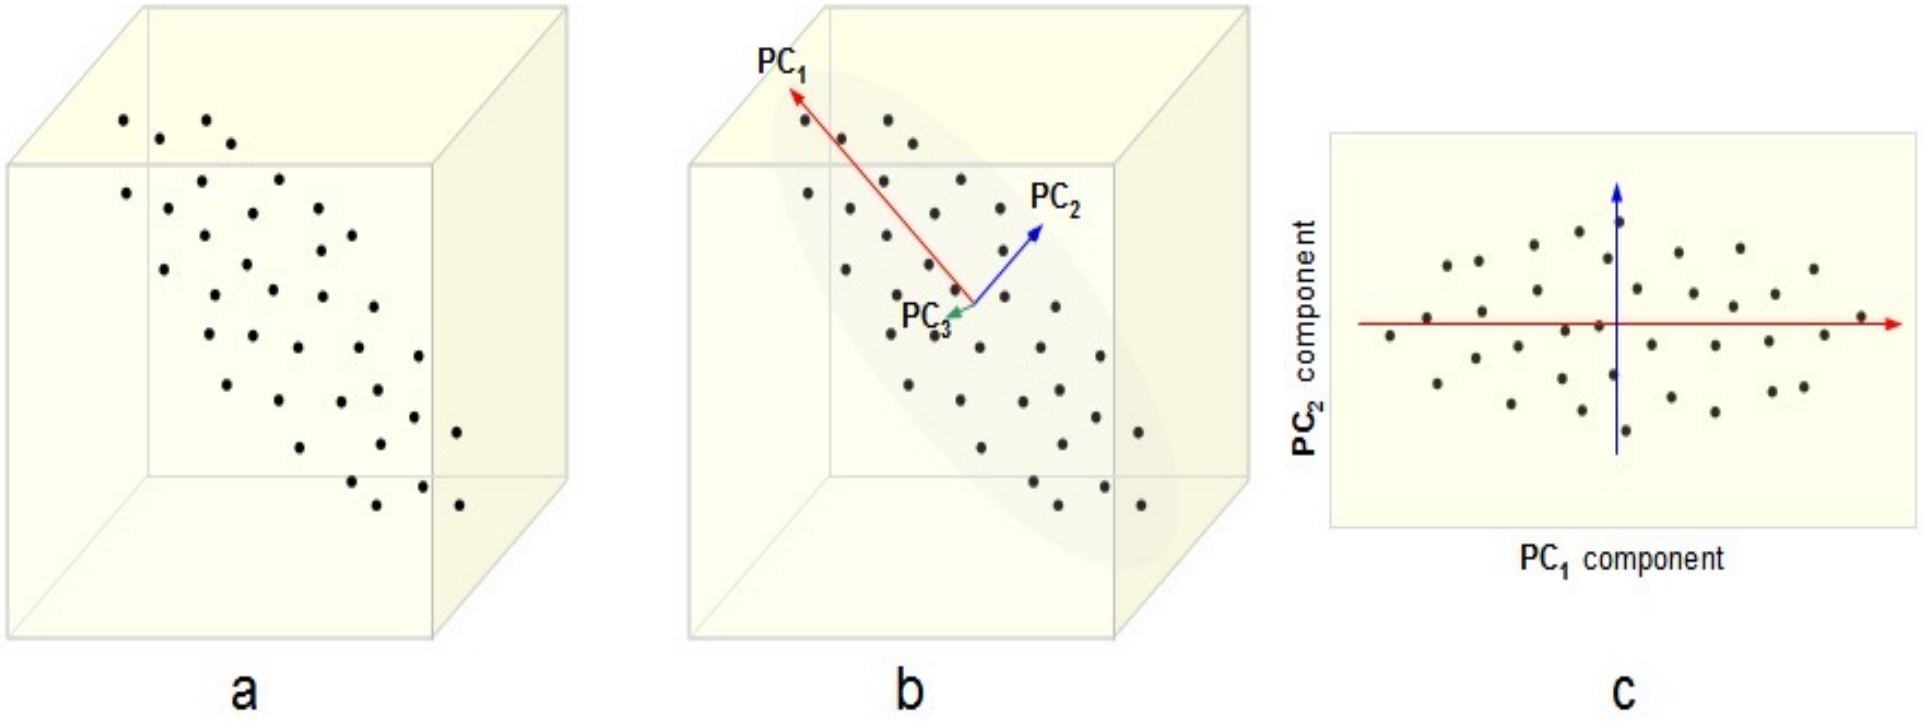

What are Principal Components?

- Some intuition:
  - Each PC is a linear combination (weighted average) of all the variables
  - Each PC is like a “created” variable, combines a little bit of every variable
  - PCs are ordered in decreasing importance: 1^{st} PC > 2^{nd} PC >…
  - PCs are no longer interpretable: they are not “real” variables

- PCA as data visualization tool:
  - Plot 1^{st} and 2^{nd} PCs on a scatter plot
  - The “best” scatter plot that captures the “most info” of data
  - Great for discover patterns in data


PEANUT BUTTER

nature
International journal of science

Altmetric: 264 Citations: 607 More detail »

Letter

Genes mirror geography within Europe

John Novembre ✉, Toby Johnson, Katarzyna Bryc, Zoltán Kutalik, Adam R. Boyko, Adam Auton, Amit Indap, Karen S. King, Sven Bergmann, Matthew R. Nelson, Matthew Stephens & Carlos D. Bustamante

Nature **456**, 98–101 (06 November 2008)
doi:10.1038/nature07331
Download Citation

Received: 30 May 2008
Accepted: 12 August 2008
Published: 31 August 2008
Addendum: 13 November 2008

- Data matrix:
  - 1,387 rows (people from Europe)
  - 197,145 columns (gene measurement)

- PCA plot:
  - reduces to 2 dimensions!
  - reveal insightful data pattern

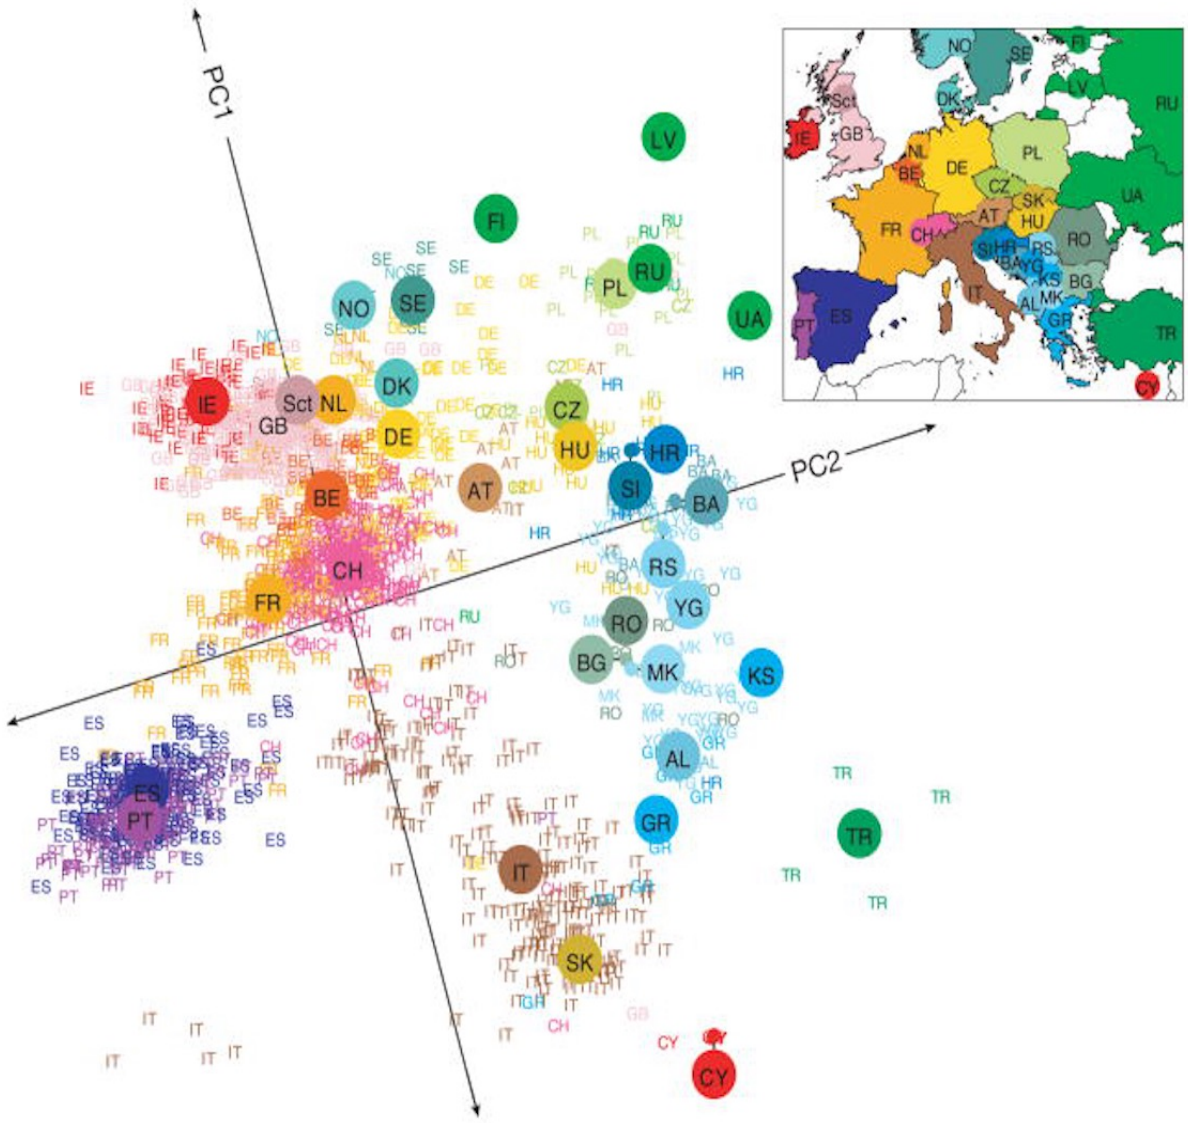

PCA in application

- Only applies to numerical variables!
  - watch out for categorical variables coded as numbers
- Can have at most p PCs for p-dimensional data!
  - mostly we just look at 1st, 2nd and maybe the 3nd PCs
- Standardize columns if variables are not comparable in scale
  - otherwise variables in large scale will dominate

Use PCA to discover data pattern

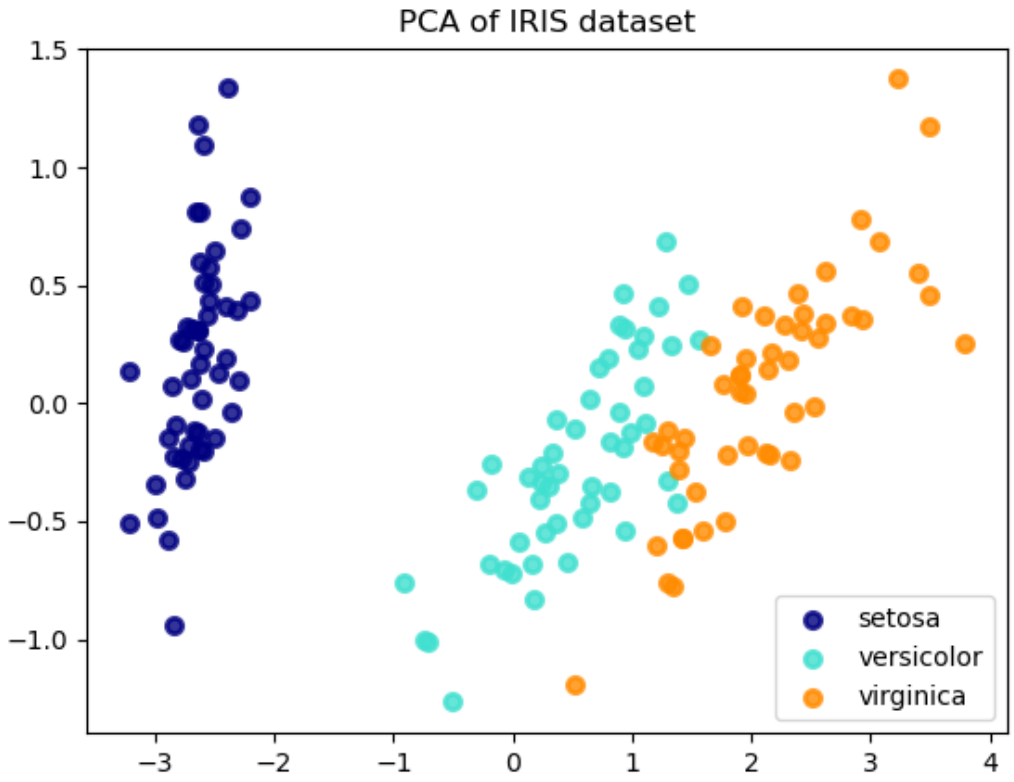

- A famous dataset of iris plants
- 4 dimensional: petal length/width, sepal length/width
- PCA scatter plot can completely separate the 3 species of iris, while any scatter plot of two variables can not

Clustering

- Finding subgroups (clusters) in a dataset.
- The observations (rows) within each cluster are similar.
- How do we define two or more observations to be similar or different?
  - use some distance metric
  - often domain-specific

Applications

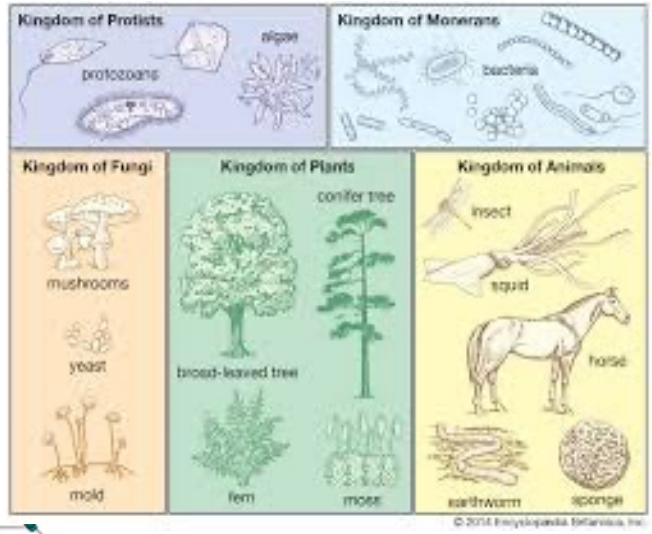

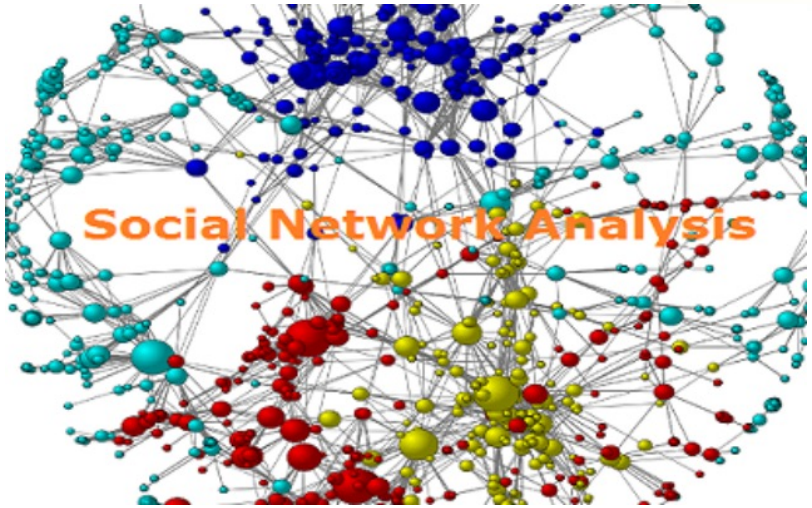

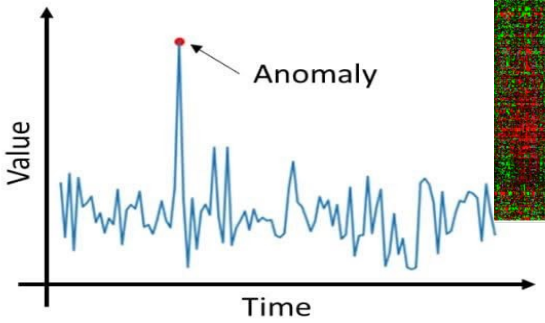

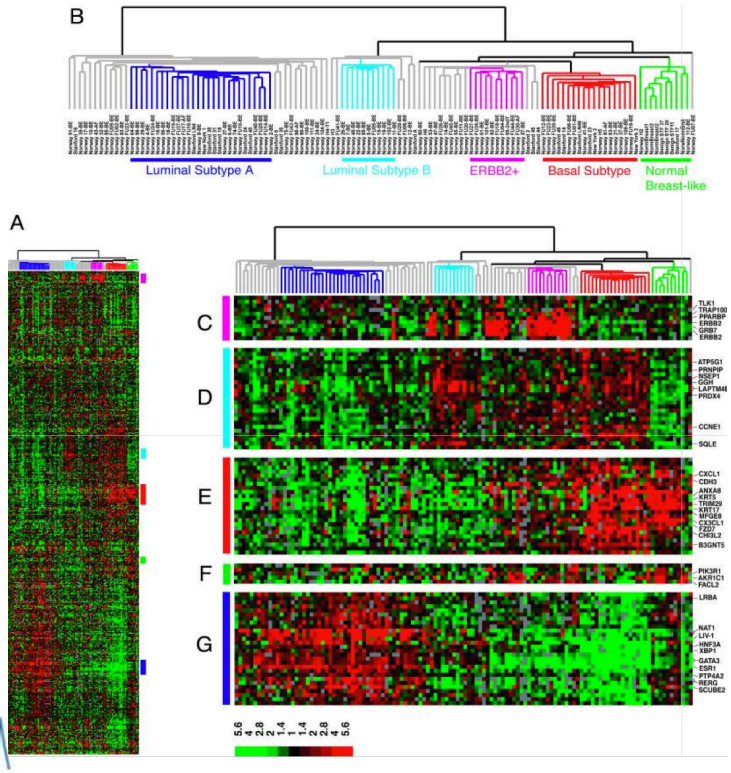

K-means

- Step 1: randomly assign cluster membership
- Step 2: iteratively update cluster centroids and membership
  - 2a: find cluster centroids
  - 2b: reassign cluster membership to the closest centroid
  - repeat until converge

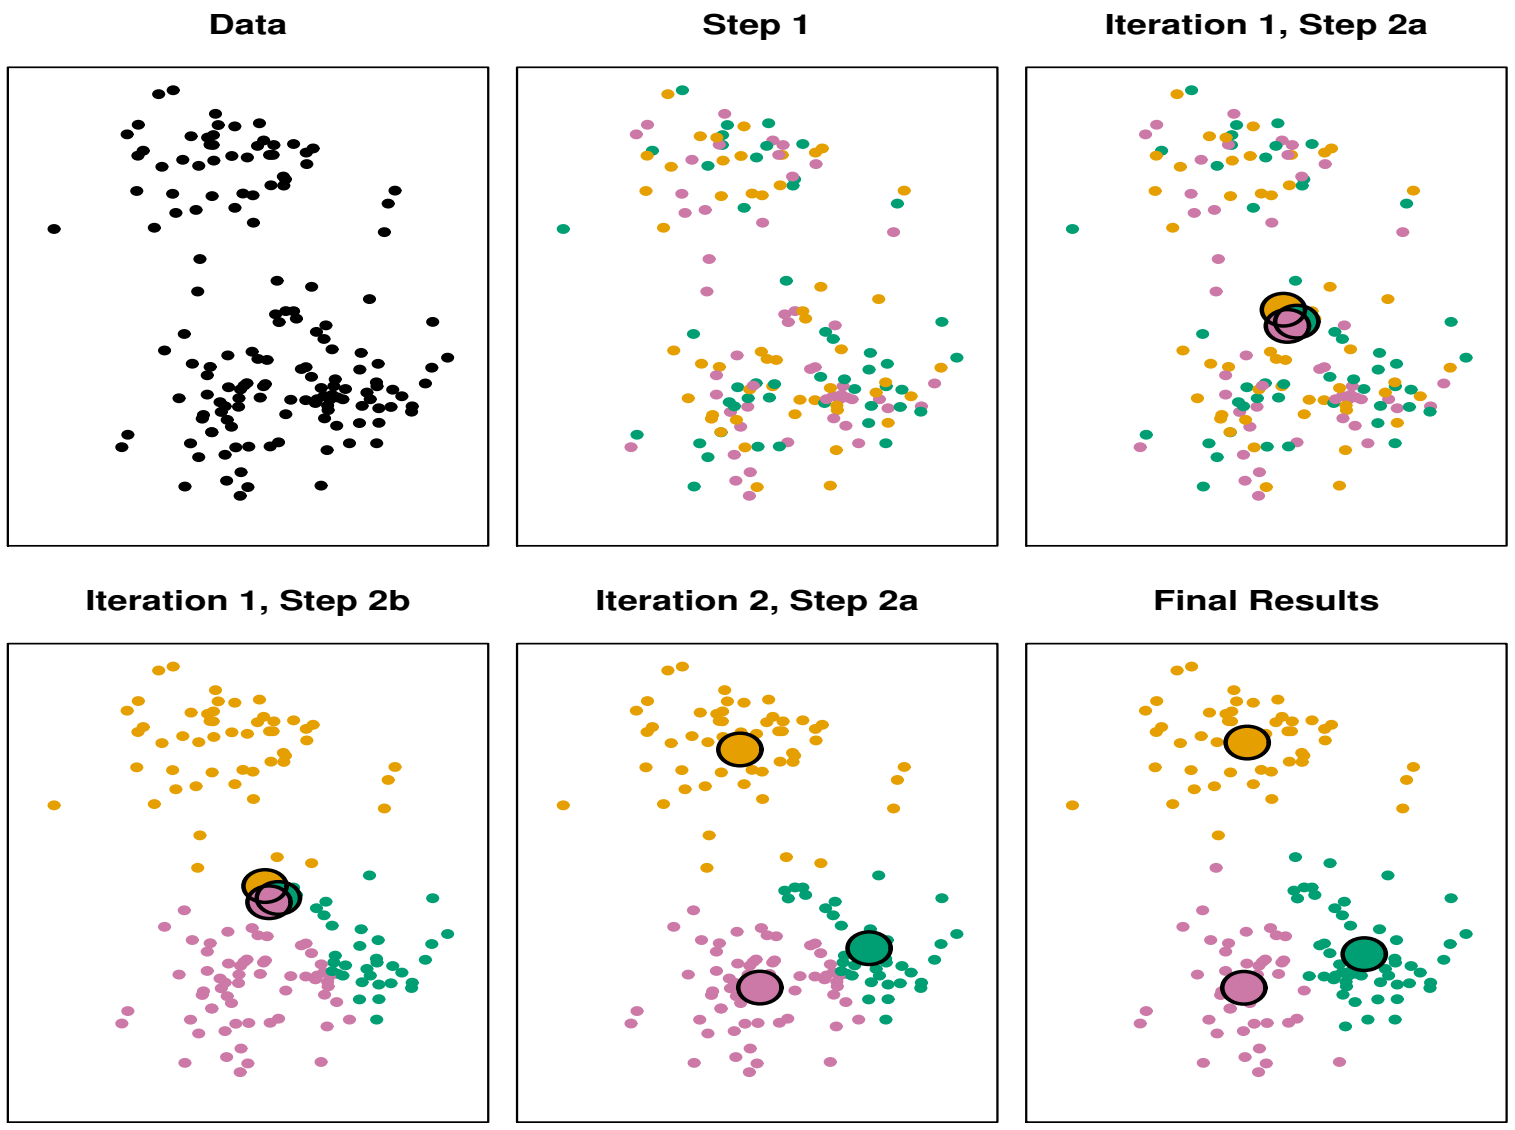

Practical issues in clustering

- Should the features be standardized?
  - for each column, subtract its mean and divide by its standard deviation
  - standardized features should have mean 0 and SD 1
- How many clusters to choose???
  - can we take a peak of the entire data?
- Robustness: how to account for noise in observations?
- Can we cluster features instead of observations?
  - or cluster both!


Congratulations!
You did it!!!